

StateSpaceDynamics.jl: A Julia package for probabilistic state space models (SSMs)

Ryan Senne^{1,2}, Zachary Loschinsky¹, Carson Loughridge¹, James Fourie¹, and Brian D. DePasquale^{1,2}

¹ Department of Biomedical Engineering, Boston University ² Graduate Program for Neuroscience, Boston University

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [Open Journals](#)

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

State-space models (SSMs) are powerful tools for modeling time series data that naturally arise across a variety of domains, including neuroscience, finance, and engineering. The unifying principle of these models is that they assume that an observation sequence, Y_1, Y_2, \dots, Y_T , is generated through an underlying hidden latent sequence, X_1, X_2, \dots, X_T . This general framework encompasses two of the most popular models for time series analysis: the Hidden Markov Model (HMM) and the (Gaussian) Linear Dynamical System (LDS, i.e., the Kalman filter). Thus, SSMs provide a probabilistic framework for describing the temporal evolution of many phenomena, and their generality naturally leads to widely applicable use cases. We introduce `StateSpaceDynamics.jl`, an open source, modular package designed to be fast, readable, and self contained for the express purpose of fitting a plurality of SSMs, easily in Julia. (Senne et al., 2024)

Statement of need

Advancements in systems neuroscience have enabled the collection of massive, multivariate, and complex time series datasets, where simultaneous recordings from hundreds to thousands of neurons are increasingly common. Interpreting these high-dimensional recordings presents a significant challenge. Recent modeling approaches suggest that neural activities can be effectively characterized by a set of latent factors evolving over a low-rank manifold. Consequently, there is a growing need for models that combine dimensionality reduction with temporal dynamics, for which state-space models (SSMs) provide a natural framework.

While advanced SSM implementations exist in Python—such as the SSM package (S. Linderman, 2022) and Dynamax (Chang et al., 2024), the Julia programming language lacks an equivalent library that meets the needs of modern neuroscientists. Existing Julia offerings, like `StateSpaceModels.jl` (Bodin, Guilherme, 2024), are limited to Gaussian observation models and rely on analytical calculation of the marginal log-likelihood function. This fundamental limitation precludes the analysis of non-Gaussian observations, which are common in neuroscience where spike counts often follow Poisson or other discrete distributions. Furthermore, the requirement for analytical computation of the marginal likelihood integral:

$$p(y_{1:T}|\theta) = \int_{x_{1:T}} p(y_{1:T}|x_{1:T}, \theta)p(x_{1:T}|\theta)dx_{1:T} \quad (1)$$

restricts the development of learning algorithms for non-conjugate observation models. To address these limitations, we have developed `StateSpaceDynamics.jl`, which provides a flexible

framework for fitting a variety of SSMs, including non-Gaussian observation models, while maintaining computational efficiency.

Package design

To address these limitations, we have developed `StateSpaceDynamics.jl`, which employs a previously advocated approach of directly maximizing the complete-data log-likelihood with respect to the hidden state path for Linear Dynamical System models (Paninski et al., 2010). By leveraging the block-tridiagonal structure of the Hessian matrix, this method allows for the exact computation of the Kalman smoother in $O(T)$ time (Paninski et al., 2010). Furthermore, it facilitates the generalization of the Rauch–Tung–Striebel (RTS) smoother to accommodate other observation noise models (e.g., Poisson and Bernoulli), requiring only the computation of the gradient and Hessian of the new model to obtain an exact maximum a posteriori (MAP) path (Macke et al., 2011).

Furthermore, given the analytical Hessian matrices are available, one can make use of this to perform an approximate EM algorithm by generating an LaPlace approximation of the posterior distribution of the latent states. One can easily make use of fast inversion algorithms of the negative Hessian (i.e., Fisher Information Matrix), which are block-tridiagonal (Rybicki & Hummer, 1990). From here one can compute the approximate second moments of the posterior distribution i.e., $\text{Cov}(X_t, X_t)$ and $\text{Cov}(X_t, X_{t-1})$, and use the analytical updates of the canonical LDS (Bishop, 2006; Paninski et al., 2010). This approach becomes exact EM in the case of Gaussian observations.

Lastly, `StateSpaceDynamics.jl` provides implementations of discrete state space models i.e., Hidden Markov Models, and the ability to fit these models using the EM algorithm. While this is not the primary development target of the package, these models are necessary for the development of hierarchical models, e.g., the switching LDS (SLDS) and the recurrent switching LDS (rSLDS) (S. W. Linderman et al., 2016; Murphy, 1998). However, the recent development of `HiddenMarkovModels.jl`, may make this feature redundant, and our future work may entail directly interfacing with this package (Dalle, 2024). Nonetheless, we provide a suite of HMM models popular in neuroscience including the classic Gaussian HMM and GLM-HMMs.

By providing these features, `StateSpaceDynamics.jl` fills a critical gap in the Julia ecosystem, offering modern computational neuroscientists the tools necessary to model complex neural data with state-space models that incorporate both dimensionality reduction and temporal dynamics.


Benchmarks

To evaluate the performance of `StateSpaceDynamics.jl`, we conducted two comprehensive benchmarking studies focusing on fitting a Gaussian Linear Dynamical System (LDS) and a Bernoulli GLM-HMM. For the Gaussian LDS benchmark, we compared our package against two alternatives: the NumPy-based Kalman Filter smoother package `pykalman` and the more recent JAX-based `Dynamax`. We intentionally excluded `StateSpaceModels.jl` from our comparison as its scope is geared towards structured time-series models.

For our Gaussian LDS experiments, we constructed a synthetic dataset as follows. The state transition matrix A was generated as a random n -dimensional rotation matrix, while the observation matrix C was created as a random $m \times n$ matrix. Both the state noise covariance Q and observation noise covariance R were set to identity matrices. To ensure a fair comparison, all packages were initialized using identical random parameters, after which we executed the EM algorithm for 100 iterations. We conducted these benchmarks using `PythonCall.jl` and `BenchmarkTools.jl` to ensure accurate timing measurements.


84 To thoroughly assess performance across different scales, we tested three sequence lengths
85 ($T = 100, 500, 1000$) and explored various dimensionality combinations, with state dimensions
86 $n = 2, 4, 8$ and observation dimensions $m = 2, 4, 8$.

Expectation-Maximization Benchmark



87 For the subsequent benchmarking experiments, we compared StateSpaceDynamics.jl, Hidden-
88 MarkovModels.jl, and Dynamax in their abilities to fit a Bernoulli HMM-GLM. We completed
89 the benchmarking on the same combinations of sequence lengths, state dimensions, and obser-
90 vation dimensions as in the Gaussian LDS benchmarking. Synthetic datasets were generated
91 for each of these combinations by sampling from a randomly generated Bernoulli HMM-GLM
92 using StateSpaceDynamics.jl. All packages were initialized to the same random parameters
93 and the EM algorithm was run for 200 iterations. We disabled the convergence criteria to
94 ensure each package completed all 200 iterations of EM. These benchmarks were conducted
95 using PythonCall.jl and BenchmarkTools.jl to ensure timing accuracy.

Bernoulli HMM-GLM Benchmark



Availability

StateSpaceDynamics.jl is publicly available under the [GNU license](https://github.com/depasquale-lab/StateSpaceDynamics.jl) at <https://github.com/depasquale-lab/StateSpaceDynamics.jl>.

Conclusion

Overall, StateSpaceDynamics.jl fills an existing gap in the Julia ecosystem for general state-space modelling that exist in Python. Importantly, our package's approach is simple enough that other candidate state-space could be easily implemented. Further, this work provides a foundation for future development of more advanced state-space models, such as the SLDS and rSLDS, which are essential for modeling complex neural data. We expect that this package will be of interest to computational neuroscientists and other researchers working with high-dimensional time series data and we are currently using its functionality in three separate projects.

Author contributions

RS was the primary developer of StateSpaceDynamics.jl, implementing the core algorithms, designing the package architecture, and writing the manuscript. ZL (Zachary Loschinsky), CL (Carson Loughridge), and JF (James Fourie) contributed to package development, including implementation of key features testing, and documentation. BDD (Brian D. DePasquale) conceived the project, provided theoretical guidance and technical oversight throughout development, secured funding, and supervised the work. All authors reviewed and approved the final manuscript.

Acknowledgements

This work was supported by the Biomedical Engineering Department at Boston University.

References

- 118
- 119 Bishop, C. M. (2006). *Pattern recognition and machine learning*. Springer.
- 120 Bodin, Guilherme, S., Raphael. (2024). *StateSpaceModels.jl*. [https://github.com/](https://github.com/LAMPSPUC/StateSpaceModels.jl)
121 [LAMPSPUC/StateSpaceModels.jl](https://github.com/LAMPSPUC/StateSpaceModels.jl)
- 122 Chang, P., Harper-Donnelly, G., Kara, A., Li, X., Linderman, S., & Murphy, K. (2024).
123 *Dynamax: State space models library in JAX* [Computer software]. Github.
- 124 Dalle, G. (2024). *HiddenMarkovModels.jl*. <https://github.com/gdalle/HiddenMarkovModels.jl>
- 125 Linderman, S. (2022). *SSM: Bayesian learning and inference for state space models*. [https://](https://github.com/lindermanlab/ssm)
126 github.com/lindermanlab/ssm
- 127 Linderman, S. W., Miller, A. C., Adams, R. P., Blei, D. M., Paninski, L., & Johnson, M. J.
128 (2016). Recurrent switching linear dynamical systems. *arXiv [Stat.ML]*.
- 129 Macke, J. H., Buesing, L., Cunningham, J. P., Yu, B. M., Shenoy, K. V., & Sahani, M. (2011).
130 *Empirical models of spiking in neural populations*. 24. [https://proceedings.neurips.cc/](https://proceedings.neurips.cc/paper_files/paper/2011/file/7143d7fbadfa4693b9eec507d9d37443-Paper.pdf)
131 [paper_files/paper/2011/file/7143d7fbadfa4693b9eec507d9d37443-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2011/file/7143d7fbadfa4693b9eec507d9d37443-Paper.pdf)
- 132 Murphy, K. P. (1998). *Switching kalman filters* [Technical report].
- 133 Paninski, L., Ahmadian, Y., Ferreira, D. G., Koyama, S., Rahnama Rad, K., Vidne, M.,
134 Vogelstein, J., & Wu, W. (2010). A new look at state-space models for neural data. *J.*
135 *Comput. Neurosci.*, 29, 107–126.
- 136 Rybicki, G. B., & Hummer, D. G. (1990). *Fast solution for the diagonal elements of the inverse*
137 *of a tridiagonal matrix*.
- 138 Senne, R., Loughridge, C., Loschinsky, Z., & DePasquale, B. (2024). *StateSpaceDynamics.jl:*
139 *A julia package for probabilistic state space models*. [https://github.com/depasquale-lab/](https://github.com/depasquale-lab/StateSpaceDynamics.jl)
140 [StateSpaceDynamics.jl](https://github.com/depasquale-lab/StateSpaceDynamics.jl). <https://github.com/depasquale-lab/StateSpaceDynamics.jl>