

# <sup>1</sup> StateSpaceDynamics.jl: A Julia package for probabilistic state space models (SSMs)

<sup>3</sup> **Ryan Senne**  <sup>1,2</sup>, **Zachary Loschinskey**  <sup>1</sup>, **James Fourie**<sup>1</sup>, **Carson Loughridge**<sup>1</sup>, and **Brian D. DePasquale**  <sup>1,2</sup>

<sup>5</sup> 1 Department of Biomedical Engineering, Boston University <sup>2</sup> Graduate Program for Neuroscience,  
<sup>6</sup> Boston University

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

## Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

## <sup>7</sup> Summary

<sup>8</sup> State-space models (SSMs) are powerful tools for modeling time series data that naturally arise in a variety of domains, including neuroscience, finance, and engineering. The unifying principle of these models is they assume an observation sequence,  $Y_1, Y_2, \dots, Y_T$ , is generated through an underlying Markovian latent sequence,  $X_1, X_2, \dots, X_T$ . This framework encompasses two popular models for time series analysis: the hidden Markov model (HMM) and the (Gaussian) linear dynamical system (LDS, i.e., the Kalman filter). Thus, SSMs provide a probabilistic framework for describing the temporal evolution of many phenomena, and their generality naturally leads to a variety of use cases. We introduce StateSpaceDynamics.jl (Senne et al., 2025), an open-source, modular package designed to be fast, readable, and self-contained for the purpose of easily fitting a plurality of SSMs in the Julia language.

## <sup>11</sup> Statement of need

<sup>12</sup> Advances in neuroscience have enabled the collection of massive, multivariate, and complex time-series datasets, where simultaneous observations from hundreds to thousands of neurons are increasingly common. Interpreting these high-dimensional datasets presents significant challenges. Recent modeling approaches suggest that neural activity can be characterized by a set of latent factors evolving within a low-dimensional manifold. Consequently, there is a growing need for models that combine dimensionality reduction with temporal dynamics, for which state-space models provide a natural framework.

<sup>13</sup> While state-space model implementations exist in Python, such as the `ssm` package (S. Linderman, 2022) and `Dynamax` (Scott W. Linderman et al., 2025), the Julia programming language lacks an equivalent that meets the needs of modern neuroscience. Existing Julia offerings, like `StateSpaceModels.jl` (Saavedra et al., 2019), can accommodate continuous-state SSMs (e.g., LDS) but are limited to Gaussian observation models and rely on analytical calculation of the marginal log-likelihood. This latter limitation precludes model inference and parameter learning for non-conjugate observations which are common in neuroscience, where neural activity follow Poisson or other discrete distributions. Packages for performing inference and learning using sampling-based methods exist in Julia (such as `Turing.jl` (Fjelde et al., 2025; Ge et al., 2018)) but are computationally inefficient compared to tailored approaches based on Expectation-Maximization (EM). For discrete SSMs, an existing Julia offering, `HiddenMarkovModels.jl` (Dalle, 2024), is efficient and scalable but not intentionally designed with the functionality for mixing models that contain both discrete and continuous latent variables, such as the switching linear dynamical system model (SLDS) (Ghahramani & Hinton, 2000; Scott W. Linderman et al., 2016) increasingly used in neuroscience. Although our primary motivation arises from challenges in modeling high-dimensional neural population

42 activity, the package is not specific to neuroscience. The algorithms and abstractions apply  
43 equally well to time-series problems in engineering, econometrics, and other fields where latent  
44 variable models and structured inference are required.

## 45 Package design

46 To address these limitations, we developed StateSpaceDynamics.jl, which provides a flexible  
47 framework for fitting a variety of SSMs—including non-Gaussian observation models and models  
48 that mix discrete and continuous latents—while maintaining computational efficiency.

49 For continuous latent-variable models, (e.g., LDS) StateSpaceDynamics.jl employs a previously  
50 advocated approach of directly maximizing the complete-data log-likelihood with respect  
51 to the hidden state path (Paninski et al., 2010). By leveraging the block tridiagonal structure  
52 of the Hessian matrix, this method allows for the exact computation of the Kalman smoother  
53 in  $\mathcal{O}(T)$  time (Paninski et al., 2010). Furthermore, it facilitates the generalization of the  
54 Rauch–Tung–Striebel (RTS) smoother to accommodate other observation models (e.g., Poisson  
55 and Bernoulli), requiring only the computation of the gradient and Hessian of the new model  
56 to obtain an exact maximum a posteriori (MAP) path (Macke et al., 2011).

57 Using analytically computable Hessians, StateSpaceDynamics.jl performs approximate EM for  
58 non-Gaussian models via Laplace approximation of the latent posterior. Speed is maintained by  
59 using fast inversion algorithms of the negative Hessian (i.e., Fisher Information Matrix), which  
60 are block tridiagonal (Rybicki & Hummer, 1990). From here StateSpaceDynamics.jl computes  
61 the approximate second moments of the posterior i.e.,  $\text{Cov}(X_t, X_t)$  and  $\text{Cov}(X_t, X_{t-1})$ , and  
62 uses the analytical updates of the canonical LDS (Bishop, 2006; Paninski et al., 2010). It is  
63 important to note that when the observations and state-evolution process are assumed to have  
64 Gaussian errors, this approach is exactly the same as using the standard Kalman Filter and  
65 RTS-Smoother, i.e., they will give the same results.

66 Lastly, StateSpaceDynamics.jl provides implementations of discrete state-space models i.e.,  
67 hidden Markov models, and the ability to fit these models using EM. While this is not the  
68 primary development target of the package, these models are necessary for the development  
69 of hierarchical models that mix discrete and continuous latents, e.g., the switching LDS  
70 (SLDS) and the recurrent switching LDS (rSLDS) (Ghahramani & Hinton, 2000; Scott W.  
71 Linderman et al., 2016; Murphy, 1998) which have become immensely popular in neuroscience  
72 and require similarly tailored computational routines for efficient inference and learning. To  
73 illustrate the functionality of StateSpaceDynamics.jl for this model class, we include an  
74 implementation of Variational Laplace EM (vLEM) (Zoltowski et al., 2020). The development  
75 of HiddenMarkovModels.jl, may make our approach to discrete model learning redundant,  
76 and future work may entail directly interfacing with this package (Dalle, 2024). Nonetheless,  
77 we provide a suite of HMM models popular in neuroscience including the classic Gaussian  
78 HMM and a variety of input-output HMMs (Bengio & Frasconi, 1994), commonly referred to  
79 as generalized linear model-HMMs (GLM-HMMs) (Ashwood et al., 2022) in neuroscience.

80 By providing these features, StateSpaceDynamics.jl fills a critical gap in the Julia ecosystem,  
81 offering modern computational neuroscientists the tools to model complex neural data with  
82 state-space models that incorporate both dimensionality reduction and temporal dynamics.

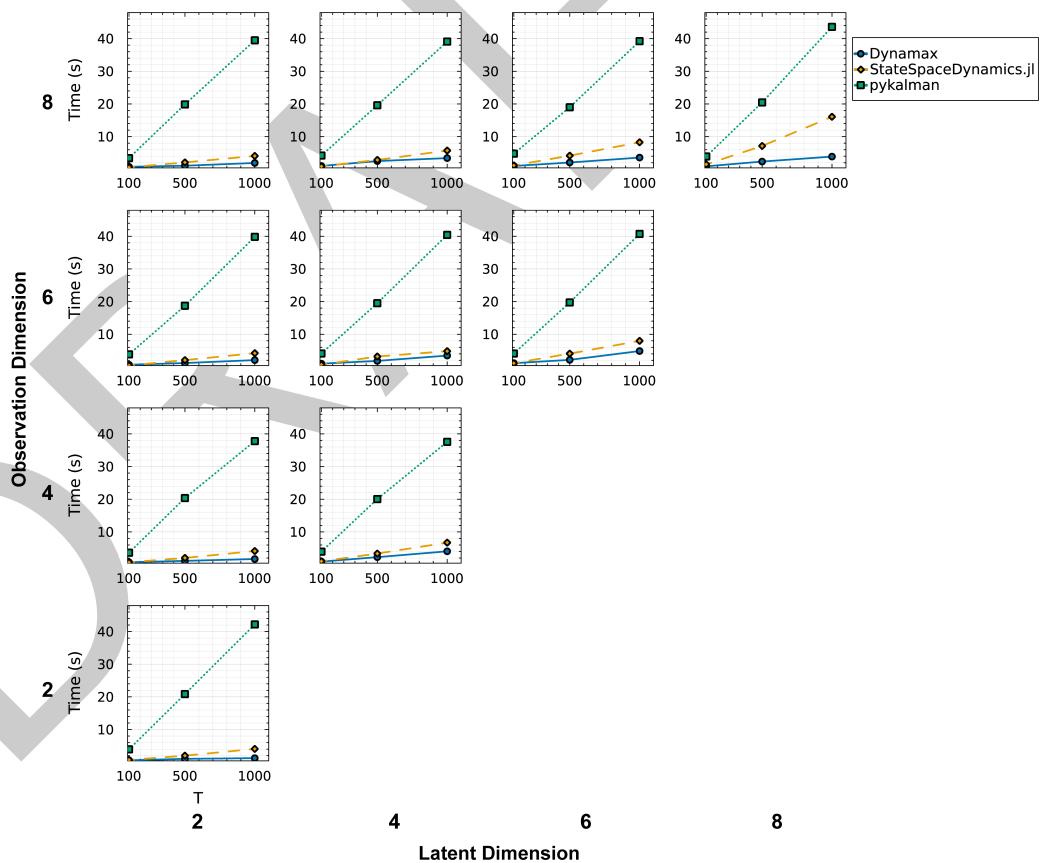
## 83 Benchmarks

84 To evaluate the performance of StateSpaceDynamics.jl, we conducted two benchmarking  
85 studies focusing on fitting a Gaussian LDS and a Gaussian HMM. For the Gaussian LDS  
86 benchmark, we compared our package against two alternatives: the NumPy-based Kalman  
87 filter-smoother package pykalman and the more recent JAX-based Dynamax. We intentionally  
88 excluded StateSpaceModels.jl from our comparison as its scope is geared towards structured

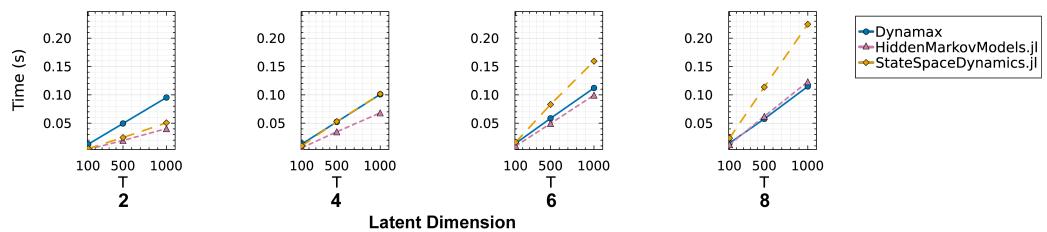
89 time-series models. Dynamax was properly JIT-compiled using the `jax.jit` function prior to  
 90 benchmarking to ensure fair comparison.

91 For our Gaussian LDS experiments, we constructed a synthetic dataset as follows. The state  
 92 transition matrix  $A$  was generated as a random  $n$ -dimensional rotation matrix, while the  
 93 observation matrix  $C$  was created as a random  $m \times n$  matrix. Both the state noise covariance  
 94  $Q$  and observation noise covariance  $R$  were set to identity matrices. To ensure a fair comparison,  
 95 all packages were initialized using identical random parameters, after which we executed the  
 96 EM algorithm for 100 iterations. We conducted these benchmarks using `PythonCall.jl`  
 97 ([Doris, 2021](#)) and `BenchmarkTools.jl` ([Chen & Revels, 2016](#)), with the assumption that  
 98 Julia-to-Python overhead is negligible for these computationally intensive operations.

99 To thoroughly assess performance across different scales, we tested three sequence lengths  
 100 ( $T = 100, 500, 1000$ ) and explored multiple dimensionality settings, with state dimensions  
 101  $n = 2, 4, 6, 8$  and observation dimensions  $m = 2, 4, 6, 8$ . In all cases, we restricted evaluations  
 102 to settings where the latent dimension was less than or equal to the observation dimension.  
 103 Finally, it is worth noting that Dynamax includes a temporally parallel smoother with  $\mathcal{O}(\log T)$   
 104 complexity. We did not include this method in our comparisons because it is GPU-specific  
 105 and incompatible with our direct optimization approach, which is designed for inference in  
 106 non-conjugate models.



107 For the second benchmarking study, we compared `StateSpaceDynamics.jl`, `HiddenMarkovModels.jl`,  
 108 and `Dynamax` in their ability to fit a Gaussian HMM. Once again, we ensured that `Dynamax`  
 109 was JIT-compiled for a fair comparison. To construct synthetic datasets, we sampled from a  
 110 Gaussian HMM with randomly selected emission models, transition matrices, and initial state  
 111 distributions. Each package was initialized using identical random parameters to maintain  
 112 consistency. EM was run for 100 iterations.



In our benchmarking, we find that for the LDS, both `StateSpaceDynamics.jl` and `Dynamax` are faster than `pykalman` across all sequence lengths and dimension configurations. More generally, `StateSpaceDynamics.jl` and `Dynamax` exhibit similar performance at lower sequence lengths (with `Dynamax` slightly outperforming `StateSpaceDynamics.jl`). However, `Dynamax` exhibits superior scaling in both the dimensions of the state and observation matrices as well as the temporal sequence length. In our current implementation, the Hessian is represented as a sparse matrix with block tridiagonal structure, resulting in  $\mathcal{O}(Tn^2)$  memory scaling — which is optimal. However, we do not yet exploit this structure fully during inference. In particular, our solver does not leverage specialized routines for block-banded systems (e.g., the block Thomas algorithm), which can result in unnecessary fill-in and degraded performance at large  $T$ . Future versions will use banded or block tridiagonal solvers to achieve truly linear-time inference.

In our HMM benchmarks, `HiddenMarkovModels.jl` outperforms both `StateSpaceDynamics.jl` and `Dynamax` across most sequence lengths and state dimensions, with `Dynamax` only becoming slightly faster for high state dimensions and long sequence lengths. `StateSpaceDynamics.jl` outperforms `Dynamax` at low state dimensions for all sequence lengths but exhibits worse scaling with the number of states, allowing `Dynamax` to overtake it as the number of states increases. These results, combined with our primary development goals in hierarchical SSMs, highlight the benefits of interfacing with `HiddenMarkovModels.jl` for HMM-specific functionality. Efforts are currently underway to make this interface seamless.

Taken together, these benchmarks demonstrate the competitiveness of `StateSpaceDynamics.jl` for fitting state-space models. Our benchmarks are available in the benchmarking folder of our repository, and instructions for running these are available in a `README.md` file.

## Availability

`StateSpaceDynamics.jl` is publicly available under the [GNU license](#) at <https://github.com/depasquale-lab/StateSpaceDynamics.jl>.

## Future Directions

The current release of `StateSpaceDynamics.jl` emphasizes efficient CPU-based implementations and analytically derived gradients and Hessians for commonly used observation models. Several avenues of future development will broaden the scope and accessibility of the package. First, we plan to add optional support for automatic differentiation (AD) using Julia's AD ecosystem (e.g., `ForwardDiff.jl` ([Revels et al., 2016](#)), `Zygote.jl` ([Innes, 2018](#)), `DifferentiationInterface.jl` ([Dalle & Hill, 2025](#))). This will allow users to prototype new observation models without requiring hand-coded derivatives, while maintaining the existing optimized implementations for speed-critical cases. Second, we aim to extend hardware support to GPU backends by exploiting Julia's GPU array abstractions and block-tridiagonal solvers, enabling large-scale inference with temporally parallel methods. Finally, we plan to expand parameter inference options beyond maximum likelihood and Laplace-EM, including Bayesian approaches via variational inference and interoperability with probabilistic programming frameworks such as `Turing.jl`. Together, these developments will further enhance the package's flexibility, performance, and utility across scientific disciplines.

## 153 Conclusion

154 StateSpaceDynamics.jl fills an existing gap in the Julia ecosystem for general state-space  
155 modeling that exists in Python. Importantly, our package's approach is simple enough that  
156 other candidate state-space models can be easily implemented. Further, this work provides a  
157 foundation for future development of more advanced state-space models, such as the rSLDS,  
158 which are essential for modeling complex neural data. We expect that this package will be of  
159 interest to computational neuroscientists and other researchers working with high-dimensional  
160 time series data and we are currently using its functionality in three separate projects.

## 161 Author contributions

162 RS (Ryan Senne) was the primary developer of StateSpaceDynamics.jl, implementing the  
163 core algorithms, designing the package architecture, and writing the manuscript. ZL (Zachary  
164 Loschinsky) was the secondary developer, whose contributions include optimizing and extending  
165 HMM/GLM-HMM functionality, implementing core multi-trial EM algorithms, and assisting  
166 with SLDS development. CL (Carson Loughridge) and JF (James Fourie) contributed to  
167 package development, including implementation of key features, testing, and documentation.  
168 BDD (Brian D. DePasquale) conceived the project, provided theoretical guidance and technical  
169 oversight throughout development, secured funding, and supervised the work. All authors  
170 reviewed and approved the final manuscript.

## 171 Acknowledgements

172 This work was supported by the Biomedical Engineering Department at Boston University.

## 173 References

- 174 Ashwood, Z. C., Roy, N. A., Stone, I. R., Urai, A. E., Churchland, A. K., Pouget, A., & Pillow,  
175 J. W. (2022). Mice alternate between discrete strategies during perceptual decision-making.  
176 *Nature Neuroscience*, 25(2), 201–212. <https://doi.org/10.1038/s41593-021-01007-z>
- 177 Bengio, Y., & Frasconi, P. (1994). An input output HMM architecture. In G. Tesauro,  
178 D. Touretzky, & T. Leen (Eds.), *Advances in neural information processing systems*  
179 (Vol. 7). MIT Press. [https://proceedings.neurips.cc/paper\\_files/paper/1994/file/8065d07da4a77621450aa84fee5656d9-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/1994/file/8065d07da4a77621450aa84fee5656d9-Paper.pdf)
- 181 Bishop, C. M. (2006). *Pattern recognition and machine learning*. Springer.
- 182 Chen, J., & Revels, J. (2016). Robust benchmarking in noisy environments. *arXiv e-Prints*.  
183 <https://arxiv.org/abs/1608.04295>
- 184 Dalle, G. (2024). HiddenMarkovModels.jl: Generic, fast and reliable state space modeling.  
185 *Journal of Open Source Software*, 9(96), 6436. <https://doi.org/10.21105/joss.06436>
- 186 Dalle, G., & Hill, A. (2025). *A common interface for automatic differentiation*. <https://arxiv.org/abs/2505.05542>
- 188 Doris, C. (2021). *PythonCall.jl*. <https://github.com/JuliaPy/PythonCall.jl>.
- 189 Fjelde, T. E., Xu, K., Widmann, D., Tarek, M., Pfiffer, C., Trapp, M., Axen, S. D., Sun,  
190 X., Hauru, M., Yong, P., Tebbutt, W., Ghahramani, Z., & Ge, H. (2025). Turing.jl: A  
191 general-purpose probabilistic programming language. *ACM Trans. Probab. Mach. Learn.*  
192 <https://doi.org/10.1145/3711897>
- 193 Ge, H., Xu, K., & Ghahramani, Z. (2018). Turing: A language for flexible probabilistic  
194 inference. In A. Storkey & F. Perez-Cruz (Eds.), *Proceedings of the twenty-first international*

- 195 conference on artificial intelligence and statistics (Vol. 84, pp. 1682–1690). PMLR.  
196 <https://proceedings.mlr.press/v84/ge18b.html>
- 197 Ghahramani, Z., & Hinton, G. E. (2000). Variational learning for switching state-space models.  
198 *Neural Computation*, 12(4), 831–864. <https://doi.org/10.1162/089976600300015619>
- 199 Innes, M. (2018). Don't unroll adjoint: Differentiating SSA-form programs. *CoRR*,  
200 [abs/1810.07951](https://arxiv.org/abs/1810.07951). <http://arxiv.org/abs/1810.07951>
- 201 Linderman, S. (2022). *SSM: Bayesian learning and inference for state space models*. <https://github.com/lindermanlab/ssm>
- 202
- 203 Linderman, Scott W., Chang, P., Harper-Donnelly, G., Kara, A., Li, X., Duran-Martin, G., &  
204 Murphy, K. (2025). Dynamax: A Python package for probabilistic state space modeling  
205 with JAX. *Journal of Open Source Software*, 10(108), 7069. <https://doi.org/10.21105/joss.07069>
- 206
- 207 Linderman, Scott W., Miller, A. C., Adams, R. P., Blei, D. M., Paninski, L., & Johnson,  
208 M. J. (2016). Recurrent switching linear dynamical systems. *arXiv [Stat.ML]*. <https://doi.org/10.48550/arXiv.1610.08466>
- 209
- 210 Macke, J. H., Buesing, L., Cunningham, J. P., Yu, B. M., Shenoy, K. V., & Sahani, M. (2011).  
211 *Empirical models of spiking in neural populations*. 24. [https://proceedings.neurips.cc/paper\\_files/paper/2011/file/7143d7fbadfa4693b9eec507d9d37443-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2011/file/7143d7fbadfa4693b9eec507d9d37443-Paper.pdf)
- 212
- 213 Murphy, K. P. (1998). *Switching Kalman filters* [Technical report]. <https://www.cs.ubc.ca/~murphyk/Papers/skf.pdf>
- 214
- 215 Paninski, L., Ahmadian, Y., Ferreira, D. G., Koyama, S., Rahnama Rad, K., Vidne, M.,  
216 Vogelstein, J., & Wu, W. (2010). A new look at state-space models for neural data. *J. Comput. Neurosci.*, 29, 107–126. <https://doi.org/10.1007/s10827-009-0179-x>
- 217
- 218 Revels, J., Lubin, M., & Papamarkou, T. (2016). Forward-mode automatic differentiation in  
219 Julia. *arXiv:1607.07892 [Cs.MS]*. <https://arxiv.org/abs/1607.07892>
- 220 Rybicki, G. B., & Hummer, D. G. (1990). *Fast solution for the diagonal elements of the inverse  
221 of a tridiagonal matrix*. <https://numerical.recipes/whp/fast/diagonal.pdf>
- 222
- 223 Saavedra, R., Bodin, G., & Souto, M. (2019). StateSpaceModels.jl: A Julia package for  
time-series analysis in a state-space framework. *arXiv Preprint arXiv:1908.01757*.
- 224
- 225 Senne, R., Loschinskey, Z., Loughridge, C., & DePasquale, B. (2025). *StateSpaceDynamics.jl:  
A Julia package for probabilistic state space models*. <https://github.com/depasquale-lab/StateSpaceDynamics.jl>. <https://github.com/depasquale-lab/StateSpaceDynamics.jl>
- 226
- 227 Zoltowski, D., Pillow, J., & Linderman, S. (2020). A general recurrent state space framework  
228 for modeling neural dynamics during decision-making. In H. D. III & A. Singh (Eds.),  
229 *Proceedings of the 37th international conference on machine learning* (Vol. 119, pp.  
11680–11691). PMLR. <https://proceedings.mlr.press/v119/zoltowski20a.html>
- 230