

¹ StateSpaceDynamics.jl: A Julia package for probabilistic state space models (SSMs)

³ **Ryan Senne**  ^{1,2}, **Zachary Loschinskey**  ¹, **James Fourie**¹, **Carson Loughridge**¹, and **Brian D. DePasquale**  ^{1,2}

⁵ 1 Department of Biomedical Engineering, Boston University ² Graduate Program for Neuroscience,
⁶ Boston University

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

⁷ Summary

⁸ State-space models (SSMs) are powerful tools for modeling time series data that naturally arise in a variety of domains, including neuroscience, finance, and engineering. The unifying principle of these models is they assume an observation sequence, Y_1, Y_2, \dots, Y_T , is generated through an underlying Markovian latent sequence, X_1, X_2, \dots, X_T . This framework encompasses two popular models for time series analysis: the hidden Markov model (HMM) and the (Gaussian) linear dynamical system (LDS, i.e., the Kalman filter). Thus, SSMs provide a probabilistic framework for describing the temporal evolution of many phenomena, and their generality naturally leads to a variety of use cases. We introduce StateSpaceDynamics.jl (Senne et al., 2025), an open-source, modular package designed to be fast, readable, and self-contained for the purpose of easily fitting a plurality of SSMs in the Julia language.

¹¹ Statement of need

¹² Advances in neuroscience have enabled the collection of massive, multivariate, and complex time-series datasets, where simultaneous observations from hundreds to thousands of neurons are increasingly common. Interpreting these high-dimensional datasets presents significant challenges. Recent modeling approaches suggest that neural activity can be characterized by a set of latent factors evolving within a low-dimensional manifold. Consequently, there is a growing need for models that combine dimensionality reduction with temporal dynamics, for which state-space models provide a natural framework.

¹³ While state-space model implementations exist in Python, such as the `ssm` package (S. Linderman, 2022) and `Dynamax` (Scott W. Linderman et al., 2025), the Julia programming language lacks an equivalent that meets the needs of modern neuroscience. Existing Julia offerings, like `StateSpaceModels.jl` (Saavedra et al., 2019), can accommodate continuous-state SSMs (e.g., LDS) but are limited to Gaussian observation models and rely on analytical calculation of the marginal log-likelihood. This latter limitation precludes model inference and parameter learning for non-conjugate observations which are common in neuroscience, where neural activity follow Poisson or other discrete distributions. Packages for performing inference and learning using sampling-based methods exist in Julia (such as `Turing.jl` (Fjelde et al., 2025; Ge et al., 2018)) but are computationally inefficient compared to tailored approaches based on Expectation-Maximization (EM). For discrete SSMs, an existing Julia offering, `HiddenMarkovModels.jl` (Dalle, 2024), is efficient and scalable but not intentionally designed with the functionality for mixing models that contain both discrete and continuous latent variables, such as the switching linear dynamical system model (SLDS) (Ghahramani & Hinton, 2000; Scott W. Linderman et al., 2016) increasingly used in neuroscience. Although our primary motivation arises from challenges in modeling high-dimensional neural population

42 activity, the package is not specific to neuroscience. The algorithms and abstractions apply
43 equally well to time-series problems in engineering, econometrics, and other fields where latent
44 variable models and structured inference are required.

45 Package design

46 To address these limitations, we developed StateSpaceDynamics.jl, which provides a flexible
47 framework for fitting a variety of SSMs—including non-Gaussian observation models and models
48 that mix discrete and continuous latents—while maintaining computational efficiency.

49 For continuous latent-variable models, (e.g., LDS) StateSpaceDynamics.jl employs a previously
50 advocated approach of directly maximizing the complete-data log-likelihood with respect
51 to the hidden state path (Paninski et al., 2010). By leveraging the block tridiagonal structure
52 of the Hessian matrix, this method allows for the exact computation of the Kalman smoother
53 in $\mathcal{O}(T)$ time (Paninski et al., 2010). Furthermore, it facilitates the generalization of the
54 Rauch–Tung–Striebel (RTS) smoother to accommodate other observation models (e.g., Poisson
55 and Bernoulli), requiring only the computation of the gradient and Hessian of the new model
56 to obtain an exact maximum a posteriori (MAP) path (Macke et al., 2011).

57 Using analytically computable Hessians, StateSpaceDynamics.jl performs approximate EM for
58 non-Gaussian models via Laplace approximation of the latent posterior. Speed is maintained by
59 using fast inversion algorithms of the negative Hessian (i.e., Fisher Information Matrix), which
60 are block tridiagonal (Rybicki & Hummer, 1990). From here StateSpaceDynamics.jl computes
61 the approximate second moments of the posterior i.e., $\text{Cov}(X_t, X_t)$ and $\text{Cov}(X_t, X_{t-1})$, and
62 uses the analytical updates of the canonical LDS (Bishop, 2006; Paninski et al., 2010). It is
63 important to note that when the observations and state-evolution process are assumed to have
64 Gaussian errors, this approach is exactly the same as using the standard Kalman Filter and
65 RTS-Smoother, i.e., they will give the same results.

66 Lastly, StateSpaceDynamics.jl provides implementations of discrete state-space models i.e.,
67 hidden Markov models, and the ability to fit these models using EM. While this is not the
68 primary development target of the package, these models are necessary for the development
69 of hierarchical models that mix discrete and continuous latents, e.g., the switching LDS
70 (SLDS) and the recurrent switching LDS (rSLDS) (Ghahramani & Hinton, 2000; Scott W.
71 Linderman et al., 2016; Murphy, 1998) which have become immensely popular in neuroscience
72 and require similarly tailored computational routines for efficient inference and learning. To
73 illustrate the functionality of StateSpaceDynamics.jl for this model class, we include an
74 implementation of Variational Laplace EM (vLEM) (Zoltowski et al., 2020). The development
75 of HiddenMarkovModels.jl, may make our approach to discrete model learning redundant,
76 and future work may entail directly interfacing with this package (Dalle, 2024). Nonetheless,
77 we provide a suite of HMM models popular in neuroscience including the classic Gaussian
78 HMM and a variety of input-output HMMs (Bengio & Frasconi, 1994), commonly referred to
79 as generalized linear model-HMMs (GLM-HMMs) (Ashwood et al., 2022) in neuroscience.

80 By providing these features, StateSpaceDynamics.jl fills a critical gap in the Julia ecosystem,
81 offering modern computational neuroscientists the tools to model complex neural data with
82 state-space models that incorporate both dimensionality reduction and temporal dynamics.

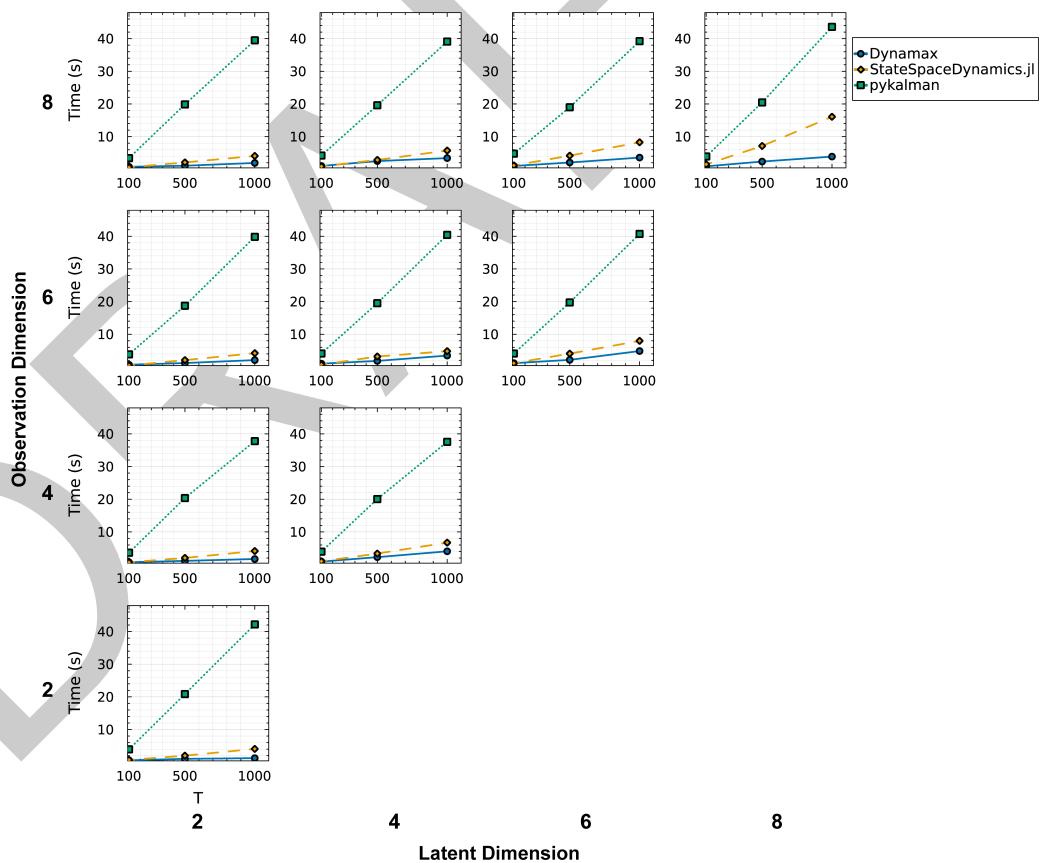
83 Benchmarks

84 To evaluate the performance of StateSpaceDynamics.jl, we conducted two benchmarking
85 studies focusing on fitting a Gaussian LDS and a Gaussian HMM. For the Gaussian LDS
86 benchmark, we compared our package against two alternatives: the NumPy-based Kalman
87 filter-smoother package pykalman and the more recent JAX-based Dynamax. We intentionally
88 excluded StateSpaceModels.jl from our comparison as its scope is geared towards structured

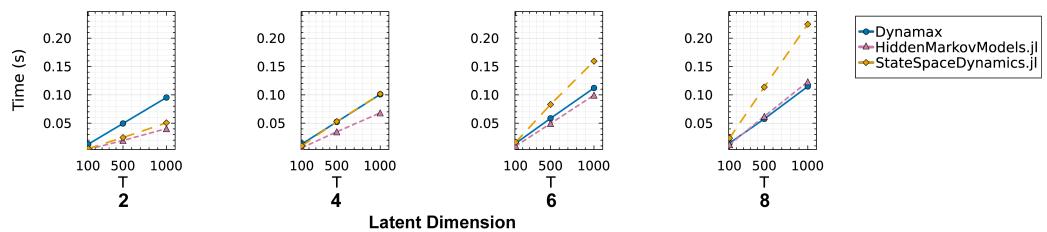
89 time-series models. Dynamax was properly JIT-compiled using the `jax.jit` function prior to
 90 benchmarking to ensure fair comparison.

91 For our Gaussian LDS experiments, we constructed a synthetic dataset as follows. The state
 92 transition matrix A was generated as a random n -dimensional rotation matrix, while the
 93 observation matrix C was created as a random $m \times n$ matrix. Both the state noise covariance
 94 Q and observation noise covariance R were set to identity matrices. To ensure a fair comparison,
 95 all packages were initialized using identical random parameters, after which we executed the
 96 EM algorithm for 100 iterations. We conducted these benchmarks using `PythonCall.jl`
 97 ([Doris, 2021](#)) and `BenchmarkTools.jl` ([Chen & Revels, 2016](#)), with the assumption that
 98 Julia-to-Python overhead is negligible for these computationally intensive operations.

99 To thoroughly assess performance across different scales, we tested three sequence lengths
 100 ($T = 100, 500, 1000$) and explored multiple dimensionality settings, with state dimensions
 101 $n = 2, 4, 6, 8$ and observation dimensions $m = 2, 4, 6, 8$. In all cases, we restricted evaluations
 102 to settings where the latent dimension was less than or equal to the observation dimension.
 103 Finally, it is worth noting that Dynamax includes a temporally parallel smoother with $\mathcal{O}(\log T)$
 104 complexity. We did not include this method in our comparisons because it is GPU-specific
 105 and incompatible with our direct optimization approach, which is designed for inference in
 106 non-conjugate models.



107 For the second benchmarking study, we compared `StateSpaceDynamics.jl`, `HiddenMarkovModels.jl`,
 108 and `Dynamax` in their ability to fit a Gaussian HMM. Once again, we ensured that `Dynamax`
 109 was JIT-compiled for a fair comparison. To construct synthetic datasets, we sampled from a
 110 Gaussian HMM with randomly selected emission models, transition matrices, and initial state
 111 distributions. Each package was initialized using identical random parameters to maintain
 112 consistency. EM was run for 100 iterations.



113 In our benchmarking, we find that for the LDS, both `StateSpaceDynamics.jl` and `Dynamax`
 114 are faster than `pykalman` across all sequence lengths and dimension configurations. More
 115 generally, `StateSpaceDynamics.jl` and `Dynamax` exhibit similar performance at lower sequence
 116 lengths (with `Dynamax` slightly outperforming `StateSpaceDynamics.jl`). However, `Dynamax`
 117 exhibits superior scaling in both the dimensions of the state and observation matrices as well
 118 as the temporal sequence length. In our current implementation, the Hessian is represented as a
 119 sparse matrix with block tridiagonal structure, resulting in $\mathcal{O}(Tn^2)$ memory scaling — which is
 120 optimal. However, we do not yet exploit this structure fully during inference. In particular, our
 121 solver does not leverage specialized routines for block-banded systems (e.g., the block Thomas
 122 algorithm), which can result in unnecessary fill-in and degraded performance at large T . Future
 123 versions will use banded or block tridiagonal solvers to achieve truly linear-time inference.
 124
 125 In our HMM benchmarks, `HiddenMarkovModels.jl` outperforms both `StateSpaceDynamics.jl`
 126 and `Dynamax` across most sequence lengths and state dimensions, with `Dynamax` only becoming
 127 slightly faster for high state dimensions and long sequence lengths. `StateSpaceDynamics.jl`
 128 outperforms `Dynamax` at low state dimensions for all sequence lengths but exhibits worse scaling
 129 with the number of states, allowing `Dynamax` to overtake it as the number of states increases.
 130 These results, combined with our primary development goals in hierarchical SSMs, highlight the
 131 benefits of interfacing with `HiddenMarkovModels.jl` for HMM-specific functionality. Efforts
 132 are currently underway to make this interface seamless.
 133
 134 Taken together, these benchmarks demonstrate the competitiveness of `StateSpaceDynamics.jl`
 135 for fitting state-space models. Our benchmarks are available in the benchmarking folder of
 136 our repository, and instructions for running these are available in a `README.md` file.

135 Availability

136 `StateSpaceDynamics.jl` is publicly available under the [GNU license](#) at <https://github.com/depasquale-lab/StateSpaceDynamics.jl>.

138 Future Directions

139 The current release of `StateSpaceDynamics.jl` emphasizes efficient CPU-based implementa-
 140 tions and analytically derived gradients and Hessians for commonly used observation
 141 models. Several avenues of future development will broaden the scope and accessibility of the
 142 package. First, we plan to add optional support for automatic differentiation (AD) using
 143 Julia's AD ecosystem (e.g., `ForwardDiff.jl` ([Revels et al., 2016](#)), `Zygote.jl` ([Innes, 2018](#)),
 144 `DifferentiationInterface.jl` ([Dalle & Hill, 2025](#))). This will allow users to prototype new
 145 observation models without requiring hand-coded derivatives, while maintaining the existing
 146 optimized implementations for speed-critical cases. Second, we aim to extend hardware support
 147 to GPU backends by exploiting Julia's GPU array abstractions and block-tridiagonal solvers,
 148 enabling large-scale inference with temporally parallel methods. Finally, we plan to expand
 149 parameter inference options beyond maximum likelihood and Laplace-EM, including Bayesian
 150 approaches via variational inference and interoperability with probabilistic programming frame-
 151 works such as `Turing.jl`. Together, these developments will further enhance the package's
 152 flexibility, performance, and utility across scientific disciplines.

153 Conclusion

154 StateSpaceDynamics.jl fills an existing gap in the Julia ecosystem for general state-space
155 modeling that exists in Python. Importantly, our package's approach is simple enough that
156 other candidate state-space models can be easily implemented. Further, this work provides a
157 foundation for future development of more advanced state-space models, such as the rSLDS,
158 which are essential for modeling complex neural data. We expect that this package will be of
159 interest to computational neuroscientists and other researchers working with high-dimensional
160 time series data and we are currently using its functionality in three separate projects.

161 Author contributions

162 RS (Ryan Senne) was the primary developer of StateSpaceDynamics.jl, implementing the
163 core algorithms, designing the package architecture, and writing the manuscript. ZL (Zachary
164 Loschinsky) was the secondary developer, whose contributions include optimizing and extending
165 HMM/GLM-HMM functionality, implementing core multi-trial EM algorithms, and assisting
166 with SLDS development. CL (Carson Loughridge) and JF (James Fourie) contributed to
167 package development, including implementation of key features, testing, and documentation.
168 BDD (Brian D. DePasquale) conceived the project, provided theoretical guidance and technical
169 oversight throughout development, secured funding, and supervised the work. All authors
170 reviewed and approved the final manuscript.

171 Acknowledgements

172 This work was supported by the Biomedical Engineering Department at Boston University.

173 References

- 174 Ashwood, Z. C., Roy, N. A., Stone, I. R., Urai, A. E., Churchland, A. K., Pouget, A., & Pillow,
175 J. W. (2022). Mice alternate between discrete strategies during perceptual decision-making.
176 *Nature Neuroscience*, 25(2), 201–212. <https://doi.org/10.1038/s41593-021-01007-z>
- 177 Bengio, Y., & Frasconi, P. (1994). An input output HMM architecture. In G. Tesauro,
178 D. Touretzky, & T. Leen (Eds.), *Advances in neural information processing systems*
179 (Vol. 7). MIT Press. https://proceedings.neurips.cc/paper_files/paper/1994/file/8065d07da4a77621450aa84fee5656d9-Paper.pdf
- 181 Bishop, C. M. (2006). *Pattern recognition and machine learning*. Springer.
- 182 Chen, J., & Revels, J. (2016). Robust benchmarking in noisy environments. *arXiv e-Prints*.
183 <https://arxiv.org/abs/1608.04295>
- 184 Dalle, G. (2024). HiddenMarkovModels.jl: Generic, fast and reliable state space modeling.
185 *Journal of Open Source Software*, 9(96), 6436. <https://doi.org/10.21105/joss.06436>
- 186 Dalle, G., & Hill, A. (2025). *A common interface for automatic differentiation*. <https://arxiv.org/abs/2505.05542>
- 188 Doris, C. (2021). *PythonCall.jl*. <https://github.com/JuliaPy/PythonCall.jl>.
- 189 Fjelde, T. E., Xu, K., Widmann, D., Tarek, M., Pfiffer, C., Trapp, M., Axen, S. D., Sun,
190 X., Hauru, M., Yong, P., Tebbutt, W., Ghahramani, Z., & Ge, H. (2025). Turing.jl: A
191 general-purpose probabilistic programming language. *ACM Trans. Probab. Mach. Learn.*
192 <https://doi.org/10.1145/3711897>
- 193 Ge, H., Xu, K., & Ghahramani, Z. (2018). Turing: A language for flexible probabilistic
194 inference. In A. Storkey & F. Perez-Cruz (Eds.), *Proceedings of the twenty-first international*

- 195 conference on artificial intelligence and statistics (Vol. 84, pp. 1682–1690). PMLR.
196 <https://proceedings.mlr.press/v84/ge18b.html>
- 197 Ghahramani, Z., & Hinton, G. E. (2000). Variational learning for switching state-space models.
198 *Neural Computation*, 12(4), 831–864. <https://doi.org/10.1162/089976600300015619>
- 199 Innes, M. (2018). Don't unroll adjoint: Differentiating SSA-form programs. *CoRR*,
200 [abs/1810.07951](https://arxiv.org/abs/1810.07951). <http://arxiv.org/abs/1810.07951>
- 201 Linderman, S. (2022). *SSM: Bayesian learning and inference for state space models*. <https://github.com/lindermanlab/ssm>
- 202
- 203 Linderman, Scott W., Chang, P., Harper-Donnelly, G., Kara, A., Li, X., Duran-Martin, G., &
204 Murphy, K. (2025). Dynamax: A Python package for probabilistic state space modeling
205 with JAX. *Journal of Open Source Software*, 10(108), 7069. <https://doi.org/10.21105/joss.07069>
- 206
- 207 Linderman, Scott W., Miller, A. C., Adams, R. P., Blei, D. M., Paninski, L., & Johnson,
208 M. J. (2016). Recurrent switching linear dynamical systems. *arXiv [Stat.ML]*. <https://doi.org/10.48550/arXiv.1610.08466>
- 209
- 210 Macke, J. H., Buesing, L., Cunningham, J. P., Yu, B. M., Shenoy, K. V., & Sahani, M. (2011).
211 *Empirical models of spiking in neural populations*. 24. https://proceedings.neurips.cc/paper_files/paper/2011/file/7143d7fbadfa4693b9eec507d9d37443-Paper.pdf
- 212
- 213 Murphy, K. P. (1998). *Switching kalman filters* [Technical report].
- 214
- 215 Paninski, L., Ahmadian, Y., Ferreira, D. G., Koyama, S., Rahnama Rad, K., Vidne, M.,
216 Vogelstein, J., & Wu, W. (2010). A new look at state-space models for neural data. *J. Comput. Neurosci.*, 29, 107–126. <https://doi.org/10.1007/s10827-009-0179-x>
- 217
- 218 Revels, J., Lubin, M., & Papamarkou, T. (2016). Forward-mode automatic differentiation in
Julia. *arXiv:1607.07892 [Cs.MS]*. <https://arxiv.org/abs/1607.07892>
- 219
- 220 Rybicki, G. B., & Hummer, D. G. (1990). *Fast solution for the diagonal elements of the inverse
of a tridiagonal matrix*.
- 221
- 222 Saavedra, R., Bodin, G., & Souto, M. (2019). StateSpaceModels.jl: A julia package for
time-series analysis in a state-space framework. *arXiv Preprint arXiv:1908.01757*.
- 223
- 224 Senne, R., Loschinsky, Z., Loughridge, C., & DePasquale, B. (2025). *StateSpaceDynamics.jl:
A julia package for probabilistic state space models*. <https://github.com/depasquale-lab/StateSpaceDynamics.jl>
- 225
- 226 Zoltowski, D., Pillow, J., & Linderman, S. (2020). A general recurrent state space framework
227 for modeling neural dynamics during decision-making. In H. D. III & A. Singh (Eds.),
228 *Proceedings of the 37th international conference on machine learning* (Vol. 119, pp.
229 11680–11691). PMLR. <https://proceedings.mlr.press/v119/zoltowski20a.html>