

¹ StateSpaceDynamics.jl: A Julia package for probabilistic state space models (SSMs)

³ **Ryan Senne**  ^{1,2}, **Zachary Loschinskey**  ¹, **Carson Loughridge**¹, **James Fourie**¹, and **Brian D. DePasquale**  ^{1,2}

⁵ Department of Biomedical Engineering, Boston University ² Graduate Program for Neuroscience,
⁶ Boston University

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Open Journals](#) 

Reviewers:

- [@openjournals](#)

Submitted: 01 January 1970

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

⁷ Summary

⁸ State-space models (SSMs) are powerful tools for modeling time series data that naturally arise
⁹ in a variety of domains, including neuroscience, finance, and engineering. The unifying principle
¹⁰ of these models is they assume an observation sequence, Y_1, Y_2, \dots, Y_T , is generated through
¹¹ an underlying Markovian latent sequence, X_1, X_2, \dots, X_T . This framework encompasses two
¹² popular models for time series analysis: the hidden Markov model (HMM) and the (Gaussian)
¹³ linear dynamical system (LDS, i.e., the Kalman filter). Thus, SSMs provide a probabilistic
¹⁴ framework for describing the temporal evolution of many phenomena, and their generality
¹⁵ naturally leads to a variety of use cases. We introduce StateSpaceDynamics.jl ([Senne et al., 2025](#)),
¹⁶ an open-source, modular package designed to be fast, readable, and self-contained for
¹⁷ the purpose of easily fitting a plurality of SSMs in the Julia language.

¹⁸ Statement of need

¹⁹ Advances in neuroscience have enabled the collection of massive, multivariate, and complex
²⁰ time-series datasets, where simultaneous observations from hundreds to thousands of neurons
²¹ are increasingly common. Interpreting these high-dimensional datasets presents significant
²² challenges. Recent modeling approaches suggest that neural activity can be characterized by
²³ a set of latent factors evolving within a low-dimensional manifold. Consequently, there is a
²⁴ growing need for models that combine dimensionality reduction with temporal dynamics, for
²⁵ which state-space models provide a natural framework.

²⁶ While state-space model implementations exist in Python, such as the SSM package ([S. Linderman, 2022](#)) and Dynamax ([Scott W. Linderman et al., 2025](#)), the Julia programming
²⁷ language lacks an equivalent that meets the needs of modern neuroscience. Existing Julia
²⁸ offerings, like StateSpaceModels.jl ([Saavedra et al., 2019](#)), can accommodate continuous-
²⁹ state SSMs (e.g., LDS) but are limited to Gaussian observation models and rely on analytical
³⁰ calculation of the marginal log-likelihood. This latter limitation precludes model inference and
³¹ parameter learning for non-conjugate observations which are common in neuroscience, where
³² neural activity follow Poisson or other discrete distributions. Packages for performing inference
³³ and learning using sampling-based methods exist in Julia (such as Turing.jl ([Fjelde et al., 2025; Ge et al., 2018](#))) but are computationally inefficient compared to tailored approaches
³⁴ based on Expectation-Maximization (EM). For discrete SSMs, an existing Julia offering,
³⁵ HiddenMarkovModels.jl ([Dalle, 2024](#)), is efficient and scalable but not intentionally designed
³⁶ with the functionality for mixing models that contain both discrete and continuous latent
³⁷ variables, such as the switching linear dynamical system model (SLDS) ([Ghahramani & Hinton, 2000; Scott W. Linderman et al., 2016](#)) increasingly used in neuroscience. Although our
³⁸ primary motivation arises from challenges in modeling high-dimensional neural population
³⁹ ⁴⁰ ⁴¹

activity, the package is not specific to neuroscience. The algorithms and abstractions apply equally well to time-series problems in engineering, econometrics, and other fields where latent variable models and structured inference are required.

45 Package design

46 To address these limitations, we developed StateSpaceDynamics.jl, which provides a flexible
47 framework for fitting a variety of SSMs—including non-Gaussian observation models and models
48 that mix discrete and continuous latents—while maintaining computational efficiency.

49 For continuous latent-variable models, (e.g., LDS) StateSpaceDynamics.jl employs a previously
50 advocated approach of directly maximizing the complete-data log-likelihood with respect
51 to the hidden state path (Paninski et al., 2010). By leveraging the block tridiagonal structure
52 of the Hessian matrix, this method allows for the exact computation of the Kalman smoother
53 in $\mathcal{O}(T)$ time (Paninski et al., 2010). Furthermore, it facilitates the generalization of the
54 Rauch–Tung–Striebel (RTS) smoother to accommodate other observation models (e.g., Poisson
55 and Bernoulli), requiring only the computation of the gradient and Hessian of the new model
56 to obtain an exact maximum a posteriori (MAP) path (Macke et al., 2011).

57 Using analytically computable Hessians, StateSpaceDynamics.jl performs approximate EM for
58 non-Gaussian models via Laplace approximation of the latent posterior. Speed is maintained by
59 using fast inversion algorithms of the negative Hessian (i.e., Fisher Information Matrix), which
60 are block tridiagonal (Rybicki & Hummer, 1990). From here StateSpaceDynamics.jl computes
61 the approximate second moments of the posterior i.e., $\text{Cov}(X_t, X_t)$ and $\text{Cov}(X_t, X_{t-1})$, and
62 uses the analytical updates of the canonical LDS (Bishop, 2006; Paninski et al., 2010). It is
63 important to note that when the observations and state-evolution process are assumed to have
64 Gaussian errors, this approach is exactly the same as using the standard Kalman Filter and
65 RTS-Smoother, i.e., they will give the same results.

66 Lastly, StateSpaceDynamics.jl provides implementations of discrete state-space models i.e.,
67 hidden Markov models, and the ability to fit these models using EM. While this is not the
68 primary development target of the package, these models are necessary for the development
69 of hierarchical models that mix discrete and continuous latents, e.g., the switching LDS
70 (SLDS) and the recurrent switching LDS (rSLDS) (Ghahramani & Hinton, 2000; Scott W.
71 Linderman et al., 2016; Murphy, 1998) which have become immensely popular in neuroscience
72 and require similarly tailored computational routines for efficient inference and learning. To
73 illustrate the functionality of StateSpaceDynamics.jl for this model class, we include an
74 implementation of the SLDS fit via structured variational EM (vEM) (Ghahramani & Hinton,
75 2000). The development of HiddenMarkovModels.jl, may make our approach to discrete
76 model learning redundant, and future work may entail directly interfacing with this package
77 (Dalle, 2024). Nonetheless, we provide a suite of HMM models popular in neuroscience
78 including the classic Gaussian HMM and a variety of input-output HMMs (Bengio & Frasconi,
79 1994), commonly referred to as generalized linear model-HMMs (GLM-HMMs) (Ashwood et
80 al., 2022) in neuroscience.

81 By providing these features, StateSpaceDynamics.jl fills a critical gap in the Julia ecosystem,
82 offering modern computational neuroscientists the tools to model complex neural data with
83 state-space models that incorporate both dimensionality reduction and temporal dynamics.

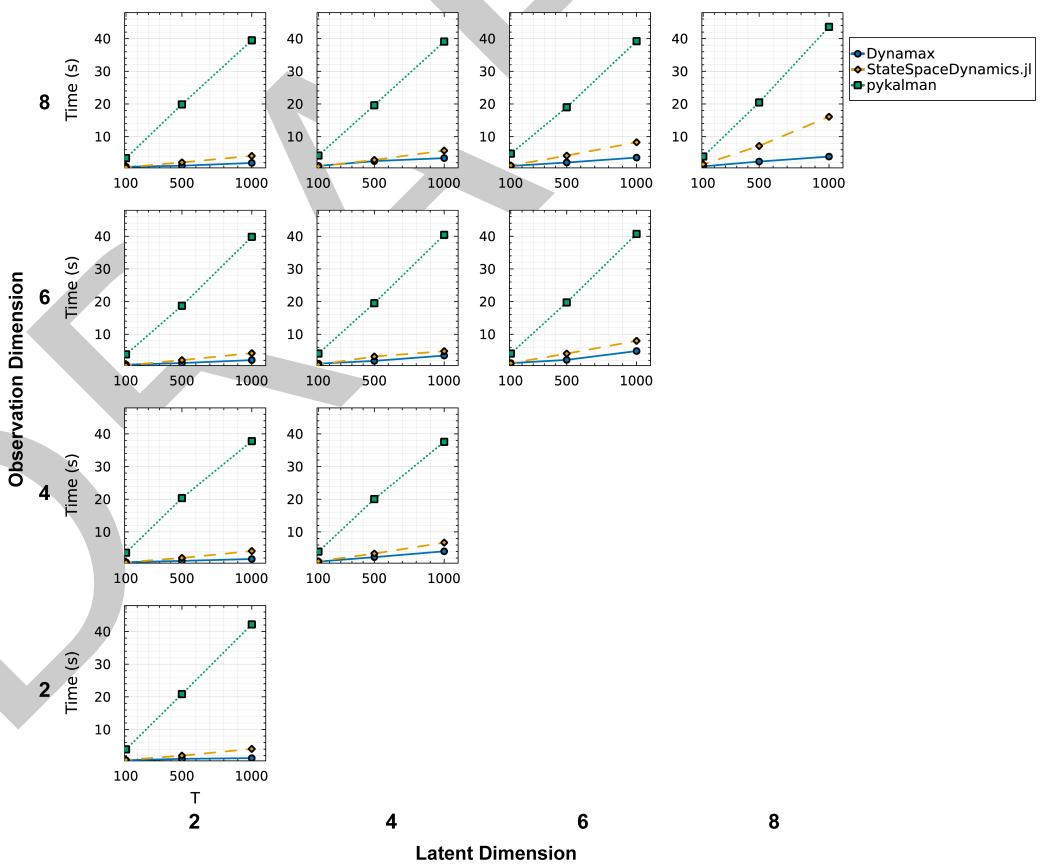
84 Benchmarks

85 To evaluate the performance of StateSpaceDynamics.jl, we conducted two benchmarking
86 studies focusing on fitting a Gaussian LDS and a Gaussian HMM. For the Gaussian LDS
87 benchmark, we compared our package against two alternatives: the NumPy-based Kalman
88 filter-smoother package pykalman and the more recent JAX-based Dynamax. We intentionally

89 excluded StateSpaceModels.jl from our comparison as its scope is geared towards structured
 90 time-series models. Dynamax was properly JIT-compiled using the `jax.jit` function prior to
 91 benchmarking to ensure fair comparison.

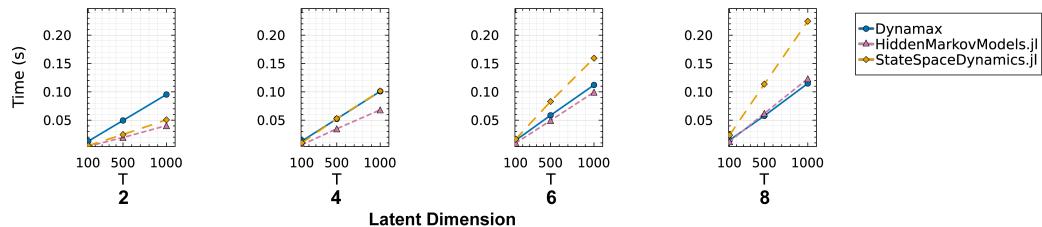
92 For our Gaussian LDS experiments, we constructed a synthetic dataset as follows. The state
 93 transition matrix A was generated as a random n -dimensional rotation matrix, while the
 94 observation matrix C was created as a random $m \times n$ matrix. Both the state noise covariance
 95 Q and observation noise covariance R were set to identity matrices. To ensure a fair comparison,
 96 all packages were initialized using identical random parameters, after which we executed the
 97 EM algorithm for 100 iterations. We conducted these benchmarks using PythonCall.jl
 98 (Doris, 2021) and BenchmarkTools.jl (Chen & Revels, 2016), with the assumption that
 99 Julia-to-Python overhead is negligible for these computationally intensive operations.

100 To thoroughly assess performance across different scales, we tested three sequence lengths
 101 ($T = 100, 500, 1000$) and explored multiple dimensionality settings, with state dimensions
 102 $n = 2, 4, 6, 8$ and observation dimensions $m = 2, 4, 6, 8$. In all cases, we restricted evaluations
 103 to settings where the latent dimension was less than or equal to the observation dimension.
 104 Finally, it is worth noting that Dynamax includes a temporally parallel smoother with $\mathcal{O}(\log T)$
 105 complexity. We did not include this method in our comparisons because it is GPU-specific
 106 and incompatible with our direct optimization approach, which is designed for inference in
 107 non-conjugate models.



108 For the second benchmarking study, we compared StateSpaceDynamics.jl, HiddenMarkovModels.jl,
 109 and Dynamax in their ability to fit a Gaussian HMM. Once again, we ensured that Dynamax
 110 was JIT-compiled for a fair comparison. To construct synthetic datasets, we sampled from a
 111 Gaussian HMM with randomly selected emission models, transition matrices, and initial state
 112 distributions. Each package was initialized using identical random parameters to maintain

¹¹³ consistency. EM was run for 100 iterations.



¹¹⁴ In our benchmarking, we find that for the LDS, both `StateSpaceDynamics.jl` and `Dynamax`
¹¹⁵ are faster than `pykalman` across all sequence lengths and dimension configurations. More
¹¹⁶ generally, `StateSpaceDynamics.jl` and `Dynamax` exhibit similar performance at lower sequence
¹¹⁷ lengths (with `Dynamax` slightly outperforming `StateSpaceDynamics.jl`). However, `Dynamax`
¹¹⁸ exhibits superior scaling in both the dimensions of the state and observation matrices as well
¹¹⁹ as the temporal sequence length. In our current implementation, the Hessian is represented as a
¹²⁰ sparse matrix with block tridiagonal structure, resulting in $\mathcal{O}(Tn^2)$ memory scaling — which is
¹²¹ optimal. However, we do not yet exploit this structure fully during inference. In particular, our
¹²² solver does not leverage specialized routines for block-banded systems (e.g., the block Thomas
¹²³ algorithm), which can result in unnecessary fill-in and degraded performance at large T . Future
¹²⁴ versions will use banded or block tridiagonal solvers to achieve truly linear-time inference.

¹²⁵ In our HMM benchmarks, `HiddenMarkovModels.jl` outperforms both `StateSpaceDynamics.jl`
¹²⁶ and `Dynamax` across most sequence lengths and state dimensions, with `Dynamax` only becoming
¹²⁷ slightly faster for high state dimensions and long sequence lengths. `StateSpaceDynamics.jl`
¹²⁸ outperforms `Dynamax` at low state dimensions for all sequence lengths but exhibits worse scaling
¹²⁹ with the number of states, allowing `Dynamax` to overtake it as the number of states increases.
¹³⁰ These results, combined with our primary development goals in hierarchical SSMs, highlight the
¹³¹ benefits of interfacing with `HiddenMarkovModels.jl` for HMM-specific functionality. Efforts
¹³² are currently underway to make this interface seamless.

¹³³ Taken together, these benchmarks demonstrate the competitiveness of `StateSpaceDynamics.jl`
¹³⁴ for fitting state-space models. Our benchmarks are available in the benchmarking folder of
¹³⁵ our repository, and instructions for running these are available in a `README.md` file.

¹³⁶ Availability

¹³⁷ `StateSpaceDynamics.jl` is publicly available under the [GNU license](#) at <https://github.com/despasquale-lab/StateSpaceDynamics.jl>.

¹³⁹ Future Directions

¹⁴⁰ The current release of `StateSpaceDynamics.jl` emphasizes efficient CPU-based implementa-
¹⁴¹ tions and analytically derived gradients and Hessians for commonly used observation mod-
¹⁴² els. Several avenues of future development will broaden the scope and accessibility of the
¹⁴³ package. First, we plan to add optional support for automatic differentiation (AD) using
¹⁴⁴ Julia's AD ecosystem (e.g., `ForwardDiff.jl` ([Revels et al., 2016](#)), `Zygote.jl` ([Innes, 2018](#)),
¹⁴⁵ `DifferentiationInterface.jl` ([Dalle & Hill, 2025](#))). This will allow users to prototype new
¹⁴⁶ observation models without requiring hand-coded derivatives, while maintaining the existing
¹⁴⁷ optimized implementations for speed-critical cases. Second, we aim to extend hardware support
¹⁴⁸ to GPU backends by exploiting Julia's GPU array abstractions and block-tridiagonal solvers,
¹⁴⁹ enabling large-scale inference with temporally parallel methods. Finally, we plan to expand
¹⁵⁰ parameter inference options beyond maximum likelihood and Laplace-EM, including Bayesian

151 approaches via variational inference and interoperability with probabilistic programming frame-
152 works such as Turing.jl. Together, these developments will further enhance the package's
153 flexibility, performance, and utility across scientific disciplines.

154 Conclusion

155 StateSpaceDynamics.jl fills an existing gap in the Julia ecosystem for general state-space
156 modeling that exists in Python. Importantly, our package's approach is simple enough that
157 other candidate state-space models can be easily implemented. Further, this work provides a
158 foundation for future development of more advanced state-space models, such as the rSLDS,
159 which are essential for modeling complex neural data. We expect that this package will be of
160 interest to computational neuroscientists and other researchers working with high-dimensional
161 time series data and we are currently using its functionality in three separate projects.

162 Author contributions

163 RS (Ryan Senne) was the primary developer of StateSpaceDynamics.jl, implementing the
164 core algorithms, designing the package architecture, and writing the manuscript. ZL (Zachary
165 Loschinsky) was the secondary developer, whose contributions include optimizing and extending
166 HMM/GLM-HMM functionality, implementing core multi-trial EM algorithms, and assisting
167 with SLDS development. CL (Carson Loughridge) and JF (James Fourie) contributed to
168 package development, including implementation of key features, testing, and documentation.
169 BDD (Brian D. DePasquale) conceived the project, provided theoretical guidance and technical
170 oversight throughout development, secured funding, and supervised the work. All authors
171 reviewed and approved the final manuscript.

172 Acknowledgements

173 This work was supported by the Biomedical Engineering Department at Boston University.

174 References

- 175 Ashwood, Z. C., Roy, N. A., Stone, I. R., Urai, A. E., Churchland, A. K., Pouget, A., & Pillow,
176 J. W. (2022). Mice alternate between discrete strategies during perceptual decision-making.
177 *Nature Neuroscience*, 25(2), 201–212. <https://doi.org/10.1038/s41593-021-01007-z>
- 178 Bengio, Y., & Frasconi, P. (1994). An input output HMM architecture. In G. Tesauro,
179 D. Touretzky, & T. Leen (Eds.), *Advances in neural information processing systems*
180 (Vol. 7). MIT Press. https://proceedings.neurips.cc/paper_files/paper/1994/file/8065d07da4a77621450aa84fee5656d9-Paper.pdf
- 182 Bishop, C. M. (2006). *Pattern recognition and machine learning*. Springer.
- 183 Chen, J., & Revels, J. (2016). Robust benchmarking in noisy environments. *arXiv e-Prints*.
184 <https://arxiv.org/abs/1608.04295>
- 185 Dalle, G. (2024). HiddenMarkovModels.jl: Generic, fast and reliable state space modeling.
186 *Journal of Open Source Software*, 9(96), 6436. <https://doi.org/10.21105/joss.06436>
- 187 Dalle, G., & Hill, A. (2025). A common interface for automatic differentiation. <https://arxiv.org/abs/2505.05542>
- 189 Doris, C. (2021). *PythonCall.jl*. <https://github.com/JuliaPy/PythonCall.jl>.
- 190 Fjelde, T. E., Xu, K., Widmann, D., Tarek, M., Pfiffer, C., Trapp, M., Axen, S. D., Sun,
191 X., Hauru, M., Yong, P., Tebbutt, W., Ghahramani, Z., & Ge, H. (2025). Turing.jl: A

- 192 general-purpose probabilistic programming language. *ACM Trans. Probab. Mach. Learn.*
193 <https://doi.org/10.1145/3711897>
- 194 Ge, H., Xu, K., & Ghahramani, Z. (2018). Turing: A language for flexible probabilistic
195 inference. In A. Storkey & F. Perez-Cruz (Eds.), *Proceedings of the twenty-first international*
196 *conference on artificial intelligence and statistics* (Vol. 84, pp. 1682–1690). PMLR.
197 <https://proceedings.mlr.press/v84/ge18b.html>
- 198 Ghahramani, Z., & Hinton, G. E. (2000). Variational learning for switching state-space models.
199 *Neural Computation*, 12(4), 831–864. <https://doi.org/10.1162/089976600300015619>
- 200 Innes, M. (2018). Don't unroll adjoint: Differentiating SSA-form programs. *CoRR*,
201 [abs/1810.07951](https://arxiv.org/abs/1810.07951). <http://arxiv.org/abs/1810.07951>
- 202 Linderman, S. (2022). *SSM: Bayesian learning and inference for state space models*. <https://github.com/lindermanlab/ssm>
- 203
- 204 Linderman, Scott W., Chang, P., Harper-Donnelly, G., Kara, A., Li, X., Duran-Martin, G., &
205 Murphy, K. (2025). Dynamax: A Python package for probabilistic state space modeling
206 with JAX. *Journal of Open Source Software*, 10(108), 7069. <https://doi.org/10.21105/joss.07069>
- 207
- 208 Linderman, Scott W., Miller, A. C., Adams, R. P., Blei, D. M., Paninski, L., & Johnson,
209 M. J. (2016). Recurrent switching linear dynamical systems. *arXiv [Stat.ML]*. <https://doi.org/10.48550/arXiv.1610.08466>
- 210
- 211 Macke, J. H., Buesing, L., Cunningham, J. P., Yu, B. M., Shenoy, K. V., & Sahani, M. (2011).
212 *Empirical models of spiking in neural populations*. 24. https://proceedings.neurips.cc/paper_files/paper/2011/file/7143d7fbadfa4693b9eec507d9d37443-Paper.pdf
- 213
- 214 Murphy, K. P. (1998). *Switching kalman filters* [Technical report].
- 215
- 216 Paninski, L., Ahmadian, Y., Ferreira, D. G., Koyama, S., Rahnama Rad, K., Vidne, M.,
217 Vogelstein, J., & Wu, W. (2010). A new look at state-space models for neural data. *J. Comput. Neurosci.*, 29, 107–126. <https://doi.org/10.1007/s10827-009-0179-x>
- 218
- 219 Revels, J., Lubin, M., & Papamarkou, T. (2016). Forward-mode automatic differentiation in
Julia. *arXiv:1607.07892 [Cs.MS]*. <https://arxiv.org/abs/1607.07892>
- 220
- 221 Rybicki, G. B., & Hummer, D. G. (1990). *Fast solution for the diagonal elements of the inverse
of a tridiagonal matrix*.
- 222
- 223 Saavedra, R., Bodin, G., & Souto, M. (2019). StateSpaceModels.jl: A julia package for
time-series analysis in a state-space framework. *arXiv Preprint arXiv:1908.01757*.
- 224
- 225 Senne, R., Loschinskey, Z., Loughridge, C., & DePasquale, B. (2025). *StateSpaceDynamics.jl:
A julia package for probabilistic state space models*. <https://github.com/depasquale-lab/StateSpaceDynamics.jl>. <https://github.com/depasquale-lab/StateSpaceDynamics.jl>
- 226